

```
np.random.seed(0)
```

Question 1.1: Implement the cost function cost/objective function:

$$J(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \left[\frac{1}{N} \sum_i \max(0, 1 - y_i * (\mathbf{w} \cdot \mathbf{x}_i + b)) \right] \quad (1)$$

```
In [2]: def compute_cost(W, X, Y, reg_strength=1000):
# TODO calculate cost function
w_norm = 0.5*np.dot(W, W)**2

offsets = np.dot(W, X.T).flatten()*Y
smax = np.clip(a=1-offsets, a_min=0.0, a_max=np.inf)
hinge_loss = reg_strength*np.sum(smax)/len(smax)

cost = w_norm+hinge_loss
return cost
```

Question 1.2: Write a method that calculate the cost gradient:

$$J(\mathbf{w}) = \frac{1}{N} \sum_i \left[\frac{1}{2} \|\mathbf{w}\|^2 + C \max(0, 1 - y_i * (\mathbf{w} \cdot \mathbf{x}_i)) \right] \quad (4)$$

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \frac{1}{N} \sum_i \begin{cases} \mathbf{w} & \text{if } \max(0, 1 - y_i * (\mathbf{w} \cdot \mathbf{x}_i)) = 0 \\ \mathbf{w} - C y_i \mathbf{x}_i & \text{otherwise} \end{cases} \quad (5)$$

```
In [3]: def calculate_cost_gradient(W, X_batch, Y_batch, reg_strength=1000):
# distance = 0
dw = np.zeros(W.shape[0])
for x, y in zip(X_batch, Y_batch):
    hinge_loss = np.max((0.0, 1.0-np.dot(W.T, x)*y))
    if hinge_loss == 0:
        dw += W
    else:
        dw += W - reg_strength*x*y
return dw/len(Y_batch)
```

Question 1.3: Write a method that performs stochastic Gradient descent as follows:

- Calculate the gradient of cost function i.e. $\nabla J(\mathbf{w})$

- Calculate the gradient of cost function i.e. $\nabla J(w)$
- Update the weights in the opposite direction to the gradient: $w = w - \alpha(\nabla J(w))$
- Repeat until convergence or until 5000 epochs are reached

```
In [4]: def sgd(data, outputs, learning_rate=0.0001, max_epochs=5000):
weights = np.random.uniform(-1.0, 1.0, size=data.shape[1])
nth = 1
prev_cost = 0.0
cost_threshold = 0.01 # in percent
# stochastic gradient descent
for epoch in range(1, max_epochs):
    # shuffle to prevent repeating update cycles
    X, Y = shuffle(data, outputs)
    for x_chunk, y_chunk in zip(np.array_split(X, 100), np.array_split(Y, 100)):
        ascent = calculate_cost_gradient(weights, x_chunk, y_chunk)
        weights -= learning_rate*ascent

    # convergence check on 2^nth epoch
    if epoch == 2**nth:
        cost = compute_cost(weights, X, Y)
        print("Epoch is:{} and Cost is: {}".format(epoch, cost))
        # stoppage criterion
        if epoch == max_epochs: # or np.abs(prev_cost - cost) < converge:
            return weights
        prev_cost = cost
        nth += 1

return weights
```

Dataset

```
In [5]: data = pd.read_csv('data_banknote_authentication.csv')

Y = data.iloc[:, -1]*2-1
X = data.iloc[:, 1:4]
X.insert(loc=len(X.columns), column='intercept', value=1)

X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.4, random_state=42)
```